

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import to_categorical
from keras.callbacks import EarlyStopping
import tensorflow
%matplotlib inline

from tensorflow.keras.preprocessing.sequence import pad_sequences

df = pd.read_csv('spam.csv', delimiter=',', encoding='latin-1')
df.head()

```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```

df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, inplace=True)
df.info()

```

```

RangeIndex: 5572 entries, 0 to 5571

```

```

Data columns (total 2 columns):

```

#	Column	Non-Null Count	Dtype
0	v1	5572 non-null	object
1	v2	5572 non-null	object

```

dtypes: object(2)

```

```

memory usage: 87.2+ KB

```

```

sns.countplot(df.v1)

```

```

plt.xlabel('Label')

```

```

plt.title('Number of ham and spam messages')

```

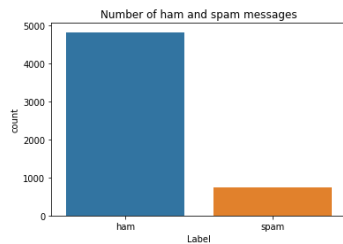
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other argumen

```

ts without an explicit keyword will result in an error or misinterpretation

```
FutureWarning
Text(0.5, 1.0, 'Number of ham and spam messages')
```



```
X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)

max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)

def RNN():
    inputs = Input(name='inputs',shape=[max_len])
    layer = Embedding(max_words,50,input_length=max_len)(inputs)
    layer = LSTM(64)(layer)
    layer = Dense(256,name='FC1')(layer)
    layer = Activation('relu')(layer)
    layer = Dropout(0.5)(layer)
    layer = Dense(1,name='out_layer')(layer)
    layer = Activation('sigmoid')(layer)
    model = Model(inputs=inputs,outputs=layer)
    return model

model = RNN()
model.summary()
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])

Model: "model"
```

Layer (type)	Output Shape	Param #
inputs (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation (Activation)	(None, 256)	0

dropout (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0

```
=====
Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0
```

```
df.columns
```

```
Index(['v1', 'v2'], dtype='object')
data=df.rename(
{
    "v1":"Category",
    "v2":"Message"
},
    axis=1
)
```

```
df.info()
```

```
RangeIndex: 5572 entries, 0 to 5571
```

```
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	v1	5572 non-null	object
1	v2	5572 non-null	object

```
dtypes: object(2)
```


```
memory usage: 87.2+ KB
```

```
data["Message Length"]=data["Message"].apply(len)
```

```
fig=plt.figure(figsize=(12,8))
```

```
sns.histplot(
    x=data["Message Length"],
    hue=data["Category"]
)
```

```
plt.title("ham & spam messege length comparision")
plt.show()
```



```
ham_desc=data[data["Category"]=="ham"]["Message Length"].describe()
```

```

spam_desc=data[data["Category"]=="spam"]["Message Length"].describe()

print("Ham Messege Length Description:\n",ham_desc)
print("*****")
print("Spam Message Length Description:\n",spam_desc)
data.describe(include="all")

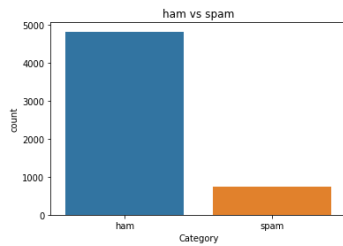
```

	Category	Message	Message Length
count	5572	5572	5572.000000
unique	2	5169	NaN
top	ham	Sorry, I'll call later	NaN
freq	4825	30	NaN
mean	NaN	NaN	80.118808
std	NaN	NaN	59.690841
min	NaN	NaN	2.000000
25%	NaN	NaN	36.000000
50%	NaN	NaN	61.000000
75%	NaN	NaN	121.000000
max	NaN	NaN	910.000000

```

data["Category"].value_counts()
ham      4825
spam      747
Name: Category, dtype: int64
sns.countplot(
    data=data,
    x="Category"
)
plt.title("ham vs spam")
plt.show()

```



```

ham_count=data["Category"].value_counts()[0]
spam_count=data["Category"].value_counts()[1]

total_count=data.shape[0]

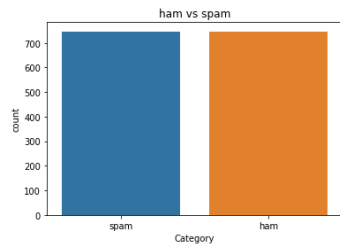
print("Ham contains:{:.2f}% of total
data.".format(ham_count/total_count*100))
print("Spam contains:{:.2f}% of total
data.".format(spam_count/total_count*100))

Ham contains:86.59% of total data.
Spam contains:13.41% of total data.
minority_len=len(data[data["Category"]=="spam"])
majority_len=len(data[data["Category"]=="ham"])
minority_indices=data[data["Category"]=="spam"].index
majority_indices=data[data["Category"]=="ham"].index
random_majority_indices=np.random.choice(
    majority_indices,
    size=minority_len,
    replace=False
)
undersampled_indices=np.concatenate([minority_indices,random_majority_indices])
df=data.loc[undersampled_indices]
df=df.sample(frac=1)

df=df.reset_index()
df=df.drop(
    columns=["index"],
)

df.shape
(1494, 3)
df["Category"].value_counts()
spam    747
ham     747
Name: Category, dtype: int64
sns.countplot(
    data=df,
    x="Category"
)
plt.title("ham vs spam")
plt.show()

```



```
df.head()
```

	Category	Message	Message Length
0	spam	Dear Voucher Holder, To claim this weeks offer...	152
1	ham	I cant pick the phone right now. Pls send a me...	51
2	spam	TheMob>Hit the link to get a premium Pink Pant...	149
3	spam	URGENT This is our 2nd attempt to contact U. Y...	136
4	spam	449050000301 You have won a £2,000 price! To ...	70

```
df["Label"]=df["Category"].map(
    {
        "ham":0,
        "spam":1
    }
)
df.head()
```

	Category	Message	Message Length	Label
0	spam	Dear Voucher Holder, To claim this weeks offer...	152	1
1	ham	I cant pick the phone right now. Pls send a me...	51	0
2	spam	TheMob>Hit the link to get a premium Pink Pant...	149	1
3	spam	URGENT This is our 2nd attempt to contact U. Y...	136	1
4	spam	449050000301 You have won a £2,000 price! To ...	70	1

```

import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

stemmer=PorterStemmer()


corpus=[]
for message in df["Message"]:
    message=re.sub("[^a-zA-Z]", " ",message)
    message=message.lower()
    message=message.split()

from tensorflow.keras.preprocessing.text import one_hot
vocab_size=10000

oneHot_doc=[one_hot(words,n=vocab_size)
             for words in corpus
            ]

df["Message Length"].describe()
count    1494.000000
mean      105.296519
std       57.379967
min        2.000000
25%       49.000000
50%      120.000000
75%      153.000000
max       450.000000
Name: Message Length, dtype: float64
fig=plt.figure(figsize=(12,8))
sns.kdeplot(
    x=df["Message Length"],
    hue=df["Category"]
)
plt.title("ham & spam messege length comparision")
plt.show()

```



```

from tensorflow.keras.preprocessing.sequence import pad_sequences
sentence_len=200
embedded_doc=pad_sequences(
    oneHot_doc,

```

```

        maxlen=sentence_len,
        padding="pre"
    )
extract_features=pd.DataFrame(
    data=embedded_doc
)
target=df["Label"]
df_final=pd.concat([extract_features,target],axis=1)
df_final.head()

```

	0	1	2	3	4	5	6	7	8	9	.	1	1	1	1	1	1	1	1	La bel	
											.	9	9	9	9	9	9	9	9		
											.	1	2	3	4	5	6	7	8	9	
0	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	1
	a	a	a	a	a	a	a	a	a	a	.	a	a	a	a	a	a	a	a	a	
	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	
1	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	0
	a	a	a	a	a	a	a	a	a	a	.	a	a	a	a	a	a	a	a	a	
	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	
2	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	1
	a	a	a	a	a	a	a	a	a	a	.	a	a	a	a	a	a	a	a	a	
	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	
3	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	1
	a	a	a	a	a	a	a	a	a	a	.	a	a	a	a	a	a	a	a	a	
	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	
4	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	1
	a	a	a	a	a	a	a	a	a	a	.	a	a	a	a	a	a	a	a	a	
	N	N	N	N	N	N	N	N	N	N	.	N	N	N	N	N	N	N	N	N	

5 rows × 201 columns

```

X=df_final.drop("Label",axis=1)
y=df_final["Label"]
from sklearn.model_selection import train_test_split
    X,
    Y,
    random_state=37,
    test_size=0.2
)
X_train,X_val,y_train,y_val=train_test_split(
    X_trainval,
    y_trainval,
    random_state=37,
    test_size=0.2
)

```



```

model = RNN()
model.summary()
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
Model: "model_1"

```

Layer (type)	Output Shape	Param #
inputs (InputLayer)	[(None, 150)]	0
embedding_1 (Embedding)	(None, 150, 50)	50000
lstm_1 (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation_2 (Activation)	(None, 256)	0
dropout_1 (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_3 (Activation)	(None, 1)	0
Total params: 96,337		
Trainable params: 96,337		
Non-trainable params: 0		